

# Demonstrating Failure of Locks (or any Concurrent System)

1. **Check Safety properties**
    - Mutual exclusion*

These must always be true.  
Two processes must not interleave certain sequences of instructions.
    - Absence of deadlock*

Deadlock is when a non-terminating system cannot respond to any signal.
  2. **Check Liveness properties**
    - Absence of starvation*
    - Fairness*
    - These must eventually be true.  
Information sent is delivered.  
That any contention must be resolved.
- **If you can demonstrate any cases in which these properties do not hold => system is not correct.**

# Example Software (not hardware) Solution to Mutual Exclusion Problem

```
/* Copyright © 2006 M. Ben-Ari. */

int wantp = 0;
int wantq = 0;

void p()
{
    while (1) {
        cout << "p non-critical section\n";
        wantp = 1;
        while (wantq == 1) {
            wantp = 0;
            wantp = 1;
        }
        cout << "p critical section\n";
        wantp = 0;
    }
}

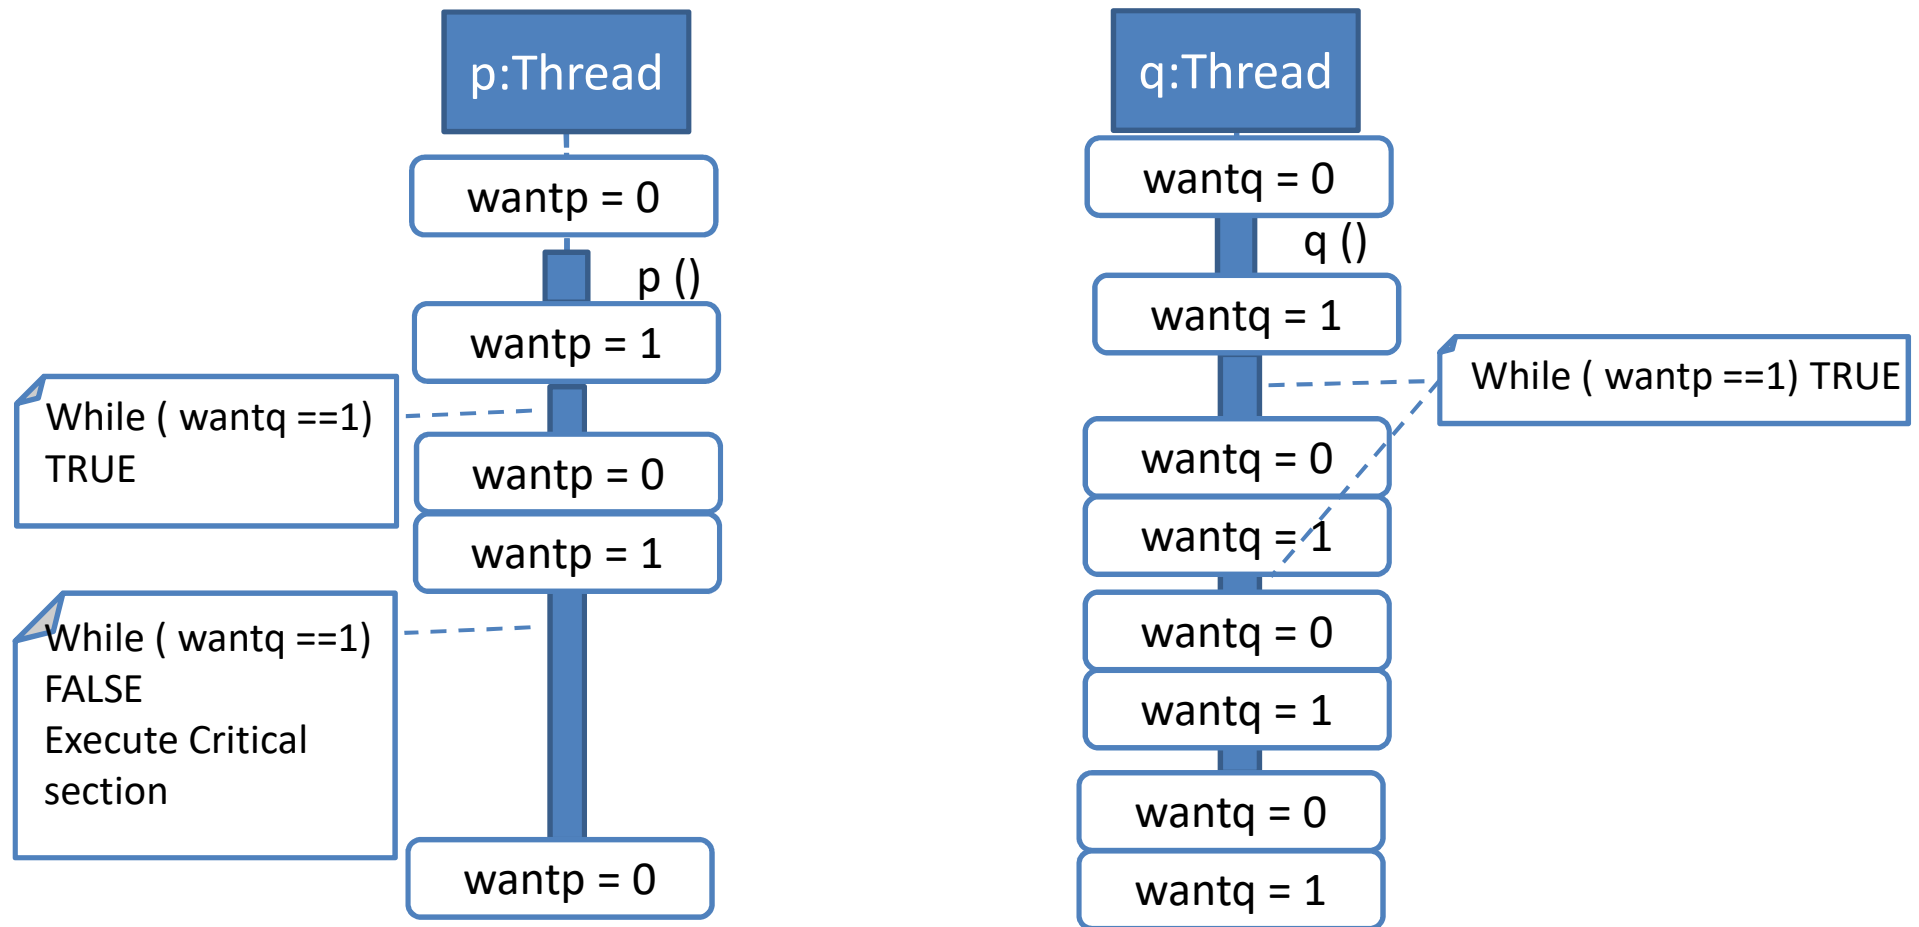
void q()
{
    while (1) {
        cout << "q non-critical section\n";
        wantq = 1;
        while (wantp == 1) {
            wantq = 0;
            wantq = 1;
        }
        cout << "q critical section\n";
        wantq = 0;
    }
}

main() {
    cobegin {
        p(); q();
    }
}
```

Exercise:

1. Turn to a partner
2. Can you find any cases where this will not work?

# Failure by Starvation



## Software Solution (cont'd)

### Proof of Failure of Software mutex Attempt:

#### 1. By Starvation

**p** sets **wantp** to 1.

**p** completes a full cycle:

Checks **wantq** Enters CS

Resets **wantp** Does non-CS

Sets **wantp** to 1

**q** sets **wantq** to 1

**q** checks **wantp**, sees **wantq**=1 & resets **wantq** to 0

**q** sets **wantq** to 1

and back



#### 2. By Livelock

**p** sets **wantp** to 1.

**p** tests **wantq**, remains in its **do** loop

**p** resets **wantp** to 0 to relinquish  
attempt to enter CS

**p** sets **wantp** to 1

**q** sets **wantq** to 1

**q** tests **wantp**, remains in its **do** loop

**q** resets **wantq** to 0 to relinquish  
attempt to enter CS

**q** sets **wantq** to 1  
etc

# Example Software (not hardware) Solution to Mutual Exclusion Problem

- This proposal has two drawbacks:

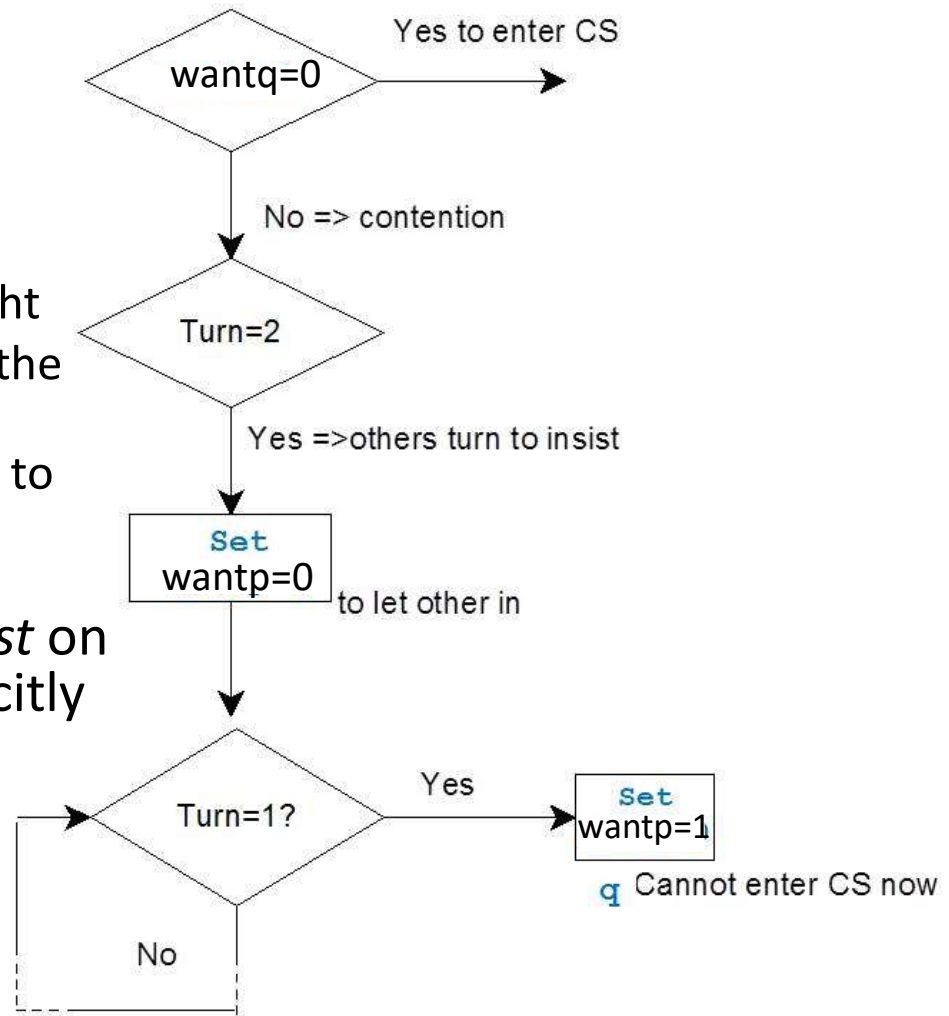
1. A process can be starved.

Can find interleavings where a process can never enter its critical section.

2. The program can *livelock* (a form of deadlock).

# Dekker's Algorithm

- A extension last proposal:
  - Add a variable to explicitly pass right to enter Critical Sections between the processes (fairness, turn-taking),
  - Last proposal had its own variable to prevent problems in absence of contention.
- In Dekker's algorithm right *to insist* on entering a Critical Section is explicitly passed between processes.
- Not safe on all hardware (optimisers can compromise this)



# Dekker's Algorithm (cont'd)

```
/* Copyright © 2006 M. Ben-Ari.
*/

int wantp = 0;
int wantq = 0;
int turn = 1; // NEW!

void p()
{
    while (1) {
        cout << "p non-CS \n";
        wantp = 1;
        while (wantq == 1) {
            wantp = 0;
            while (!(turn == 1));
            wantp = 1; }
        cout << "p CS\n";
        turn = 2;
        wantp = 0;
    }
}

void q()
{
    while (1) {
        cout << "q non-CS\n";
        wantq = 1;
        while (wantp == 1) {
            wantq = 0;
            while (!(turn == 2));
            wantq = 1; }
        cout << "q CS\n";
        turn = 1;
        wantq = 0;
    }
}

main() {
    /* As before */
}
```

# Mutual Exclusion for n Processes:

## The Bakery Algorithm

- Dekker's Algorithm solves mutual exclusion problem for 2 processes.
- Many algorithms solve  $N$  process ME problem; all are complicated and relatively slow to other methods.
- *The Bakery Algorithm* is one where processes take a numbered ticket (whose value constantly increases) when it wants to enter its CS.
- The process with the lowest current ticket gets to enter its CS.
- This algorithm is not practical because:
  - ticket numbers will be unbounded if a process is always in its critical section, and
  - even in the absence of contention it is very inefficient as each process must query the other processes for their ticket number.



```
/* Copyright (C) 2006 M. Ben-Ari. */
```

```
const int NODES = 3;
int num[NODES];
int choose[NODES];
```

```
int Max() {
int Current = 0;
int i;
for (i=0; i < NODES; i++)
    if (num[i] > Current) Current = num[i];
return Current;
}
```

```
void p(int i) {
int j;
while (1)
{
    cout << "proc " << i << " non-CS\n";
    choose[i] = 1;
    num[i] = 1 + Max();
    choose[i] = 0;
    for (j=0; j < NODES; j++)
        if (j != i)
        {
            while (choose[j]);
            while (!
                ((num[j]==0) || (num[i]<num[j])) ||
                ((num[i]==num[j]) && (i < j))) );
        }
    cout << "process " << i << " CS\n";
    num[i]=0;
}
}
```

```
main() {
int j;
for (j=0; j < NODES; j++) number[j]=0;
for (j=0; j < NODES; j++) choose[j]=0;
cobegin {
    p(0); p(1); p(2); // 3 processes here
}
}
```

## Mutual Exclusion for $N$ Processes: The Bakery Algorithm (cont'd)

## ***SECTION 2.2:* HIGHER LEVEL SUPPORT FOR MUTUAL EXCLUSION: SEMAPHORES & MONITORS**

# Example Scenario: Producer/Consumer

- Producer: creates a resource (data)
- Consumer: Uses a resource (data)
- E.g. `ps | grep "gcc" | wc`
- Don't want producers and consumers to operate in lockstep (ie atomicity)
  - Each cmd must wait for the previous output
  - Implies lots of context switching (v expensive)
- Solution (Pattern): place a fixed size buffer between producers and consumers
  - Synchronise access to buffer
  - Producer waits if buffer full; consumer waits if buffer empty

# Semaphores

- Semaphore = higher level synchronisation primitive
  - Invented by Dijkstra in 1965 as part of THE O/S project
- Implement with
  - A **counter** that is manipulated atomically via 2 operations **signal** and **wait**
  - `wait (semaphore)` : decrement, if counter is zero then block until semaphore is signalled (AKA `down()` or `P()`)
  - `signal (semaphore)` : increment counter, wake up one waiter if any (AKA `up()` or `V()`)
  - `sem_init(semaphore, counter)` : set initial counter value

# Semaphore Pseudocode

```
struct semaphore {  
    int value;  
    queue L; // list of processes  
}  
wait (S) {  
    if (s.value > 0)  
        s.value = s.value - 1;  
    else {  
        add this process to s.L;  
        block;  
    }  
}  
signal (S) {  
    if (S.L != EMPTY) {  
        remove a process P from S.L;  
        wakeup(P);  
    } else  
        s.value = s.value + 1;  
}
```

wait() / signal()  
are critical sections!  
Hence, they must be  
Executed atomically  
With respect to each  
Other.

# Blocking Semaphores

- Each semaphore has an associated queue of threads
  - When `wait()` is called by a thread
    - If semaphore is available => thread continues
    - If semaphore is unavailable, thread blocks, waits on queue
  - `signal()` opens the semaphore
    - If threads are waiting on a queue, one thread is unblocked
    - If no threads are on the queue, the signal is remembered for the next time `wait()` is called
  - NB Blocking threads are not spinning, they release the CPU to do other work

# Semaphore Initialisation

- If semaphore initialised to 1
  - First call to wait goes through
    - Semaphore value goes from 1 to 0
  - Second call to wait() blocks
    - Semaphore value stays at zero, thread goes on queue
  - If first thread calls signal()
    - Semaphore value stays at 0
    - Wakes up second thread

⇒ Acts like a mutex lock

⇒ Can use semaphores to implement locks

This is called a **binary semaphore**

# What happens if we initialise to 2?

```
struct semaphore {
    int value;
    queue L; // list of processes
}
wait (S) {
    if (s.value > 0)
        s.value = s.value -1;
    else {
        add this process to
        s.L;
        block;
    }
}
signal (S) {
    if (S.L != EMPTY){
        remove a process P
        from S.L;
        wakeup(P);
    } else
        s.value = s.value + 1;
}
```

Sem\_init(sem, 2)

Consider multiple threads:

Thread1: wait(sem)

Thread2: wait(sem)

Thread2: wait(sem) –blocks

Observations:

Initial value of semaphore =  
number of threads that can be  
active at once



# Uses of Semaphores

- Allocating a number of resources
  - Shared buffers: each time you want to access a buffer, call `wait()` => you are queued if there is no buffer available
  - Devices
- Counter is initialised to  $N$  = number of resources
- Called a **counting semaphore**
- Useful for conditional synchronisation
  - I.e. one thread is waiting for another thread to finish a piece of work before it continues

# Semaphores for Mutual Exclusion

- With semaphores, guaranteeing mutual exclusion for  $N$  processes is trivial

```
semaphore mutex = 1;

void P (int i) {
    while (1) {
        // Non Critical Section Bit
        wait(mutex) // grab the mutual exclusion semaphore
        // Do the Critical Section Bit
        signal(mutex) //grab the mutual exclusion semaphore
    }
}


int main ( ) {
    cobegin {
        P(1); P(2);
    }
}
```

# Example bounded buffer problem

- AKA producer/consumer problem
  - Buffer in memory
    - Finite size of N entries
  - A producer process inserts an entry into it
  - A consumer process removes an entry from it
- Processes are concurrent
  - => Must use a synchronisation mechanism to control access to shared variables describing buffer state

# Producer/consumer single buffer

- Simplest case
  - Single producer thread, single consumer thread
  - Single shared buffer between producer and consumer
- Requirements
  - Consumer must wait for producer to fill buffer
  - Producer must wait for consumer to empty buffer (if filled)



# Exercise:

- How many shared data values do you have?  
3 : lock, emptyBuffer flag, sharedBuffer flag
- How many semaphores do we need?
- How should we initialise them?

## Producer

```
while(1) {  
    wait(&emptyBuffer);  
    wait(&lock);  
  
    fill(&buffer);  
  
    signal(&lock);  
    signal(&fullBuffer);  
}
```

## Consumer

```
while(1) {  
    wait(&fullBuffer);  
    wait(&lock);  
  
    use(&buffer);  
  
    signal(&lock);  
    signal(&emptyBuffer);  
}
```

# Types of Semaphores

- Defined above is a general semaphore. A *binary semaphore* is a semaphore that can only take the values 0 and 1.
- Choice of which suspended process to wake gives the following definitions:
  - *Blocked-set semaphore*                      Wakes any one suspended process
  - *Blocked-queue semaphore*                      Suspended processes are kept in FIFO & woken in order of suspension
  - *Busy-wait semaphore*                      semaphore value is tested in a busy-wait loop, with atomic test. Some loop cycles may be interleaved.

# Semaphores can be hard to Use

- Complex patterns of resource usage
  - Cannot capture relationships with semaphores alone
  - Need extra state variables to record information (see Sleeping Barber later)
  - Often use semaphores such that
    - One is for mutex around state variables
    - One for each class of waiting

⇒ Produce buggy code that is hard to write

- If one coder forgets to do **V()** / **signal()** after critical section, the whole system can deadlock

# Monitors

- Need a higher level construct that groups the responsibility for correctness.
  - Supports controlled access to shared data
    - Synchronisation code added by compiler, enforced at runtime
- *Monitors* do this. They're an extension of the monolithic monitor used in OS to allocate memory etc.
  - *Encapsulate*
    - Shared data structures
    - Procedures that operate on shared data
    - Synchronization between concurrent processes that invoke these procedures
  - Ensure only one process execute a monitor procedure at once ( $\Rightarrow$ ME).
  - Guarantees only way to access the shared data is through procedures.
- Native language support in Java