

LECTURE 3: CONCURRENT & DISTRIBUTED ARCHITECTURES

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

1

Contents

- Introduction
- Flynn's Taxonomy
- MIMD: Tight & Loose Coupling
- Software Architectures for Distributed Systems:
 - Layered, Object-/Event-based, Shared Dataspace Architectures
- System Architectures
 - Centralized Architectures:
 - 2 & multi-tiered architectures
 - Fat & Thin Clients
 - Decentralized Architectures
 - Structured P2P Systems: Chord & Pastry Routing algorithms
 - Unstructured P2P Systems
 - Hybrid Systems
- Architectures V Middleware

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

2

Intro to Architectures in Concurrent & Dist'd Systems: S/w V System Architectures

- Organizing concurrent & distributed systems is mostly about the software components making up the system.
- These *software architectures* (aka *Programming Models*) dictate the organization & interaction of the various s/w components .
- The actual realization of a system requires instantiating and placing software components on real machines.
- There are many different choices that can be made in doing so.
- The final instantiation of a software architecture is referred to as a *system architecture* (aka *Machine Model*).

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

3

SECTION 3.1: CONCURRENT ARCHITECTURES & PROGRAMMING MODELS

Aside on Writing Concurrent Code

1. Identify concurrency in task
 - Do this on a piece of paper
2. Expose the concurrency when writing the task
 - Choose a *programming model* and language that allow you to express this concurrency
3. Exploit the concurrency
 - Carefully choose a language & hardware that facilitate advantage to be taken of the concurrency (often one \Leftrightarrow another)
- Value of a programming model is judged on
 - *Generality*: how well a range of different problems can be expressed for a variety of different architectures,
 - *Performance*: how efficiently compiled programs can execute on these architectures.

Lecture 3: Concurrent & Distributed Architectures

CA4006 Lecture Notes (Martin Crane 2017)

5

Parallel Programming Model

- *Definition*: Programming model comprises languages & libraries that create an abstract view of the machine.
 - Control
 - What orderings exist between operations?
 - How do different threads of control synchronize?
 - Data
 - What data is private vs. shared?
 - How is logically shared data accessed or communicated?
 - Synchronization
 - What operations can be used to coordinate parallelism?
 - What are the atomic (indivisible) operations?

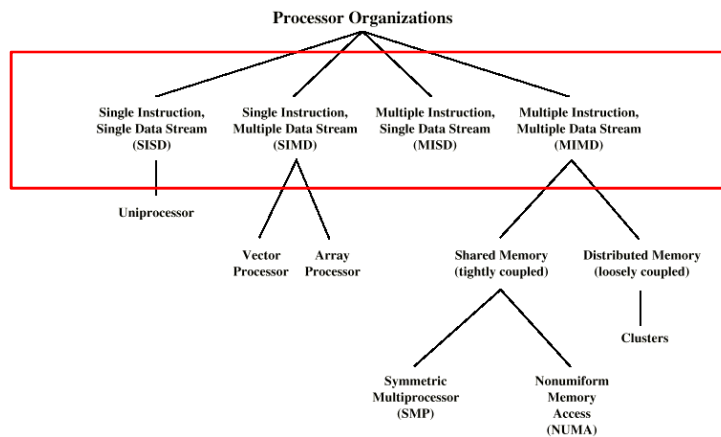
Lecture 3: Concurrent & Distributed Architectures

CA4006 Lecture Notes (Martin Crane 2017)

6

Concurrent Architecture Taxonomies

- As seen above, Michael Flynn in 1966 classified machines into a taxonomy by the number of instruction and data streams
- We examine these from standpoint of concurrent architectures



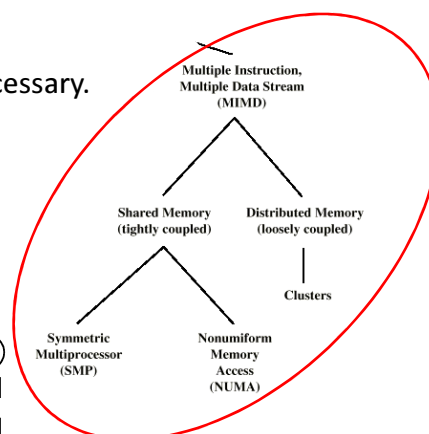

Lecture 3: Concurrent & Distributed Architectures

CA4006 Lecture Notes (Martin Crane 2017)

7

More on MIMD

- **MIMD**
- General purpose processor
- Each can process all instructions necessary.
- Further classified by method of processor communication:
 - **Tight Coupling**
 - **Loose Coupling**




Lecture 3: Concurrent & Distributed Architectures

CA4006 Lecture Notes (Martin Crane 2017)

8

Concurrent Architectures

- *Machine Model #1: Shared Memory*
- Processors all connected to a large shared memory
 - Typically Symmetric Multiprocessors (SMPs e.g. IBM SMPs)
 - Multicore chips, except caches are often shared in multicores
 - But
 - Bus is a bottleneck (interconnect performance not scalable)
 - Also, shared memory can give issues with *race conditions*
 - Can be fixed by adding locks of some sort, at performance cost




Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2015)

9

Programming Models

- *Programming Model # 1: Shared Memory*
- Program is a collection of threads of control.
- Each thread has set of private variables, e.g., local stack variables & set of shared variables, e.g., static variables
 - Implicit comms between threads writing/reading shared variables
 - Threads coordinate by synchronizing on shared variables
 - Here model used by threads calculating the sum (S) of an array



Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)


10

Concurrent Architectures (/2)

- *Machine Model #2: Distributed Memory*

Processors have own memory but typically fast interconnect

- Each processor has its own memory and cache but cannot directly access another processor's memory.
- Each “node” has a Network Interface (NI) for all communication and synchronization.
- Example: IBM SP2, Beowolf Cluster



Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)


11

Programming Models (/2)

- *Programming Model # 2: Message Passing*

Program consists of a collection of named processes.

- Usually fixed at program startup time
- Thread of control plus local address space—NO shared data.
- Logically shared data is partitioned over local processes.
- Here, similar calculation as last time.



Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)


12

Concurrent Architectures (/3)

- *Machine Model #3: Clusters*

Used for computation-intensive purposes, (Vs for IO operations such as web service or DBs.)

- Emerged as result of trends e.g. availability of low-cost cores, high speed networks & s/w for high-performance distributed computing.
- Wide applicability from small biz clusters to fastest supercomputers
- Applications that can be done however, are nonetheless limited, since s/w needs to be purpose-built per task.



Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

13

Programming Models (/3)

- *Programming Model # 3: Hybrids*

Need to run "same/similar computation" on many nodes very fast

- Common model: Hybrid MPI + OpenMP
 - Each SMP node = 1 MPI process, w MPI comm on node interconnect
 - OpenMP inside of each SMP node
- Maybe gives the highest performance?
 - **Advantage:** Could be good for heavyweight comms between nodes & lightweight threads within a node
 - **Disadvantages:**
 - Very difficult to start with OpenMP and modify for MPI
 - Very difficult to program, debug, modify and maintain
 - Generally, cannot do MPI calls within OpenMP parallel regions
 - Only people experienced in both should use this mixed prog model

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

14

SECTION 3.2: ARCHITECTURES FOR DISTRIBUTED SYSTEMS


Architectures for *Distributed Systems*

- *Introduction*
- Examine traditional centralized *distributed systems* architectures where 1 server implements most s/w components (thus functionality)
- Remote clients access the server using simple communication means.
- Also consider decentralized architectures in which machines more or less play equal roles, as well as hybrid organizations.
- From Lecture 1, one aim of distributed systems is separating applications from underlying platforms by providing a m/w layer.
- Adopting such a layer is an important architectural decision, and its main purpose is to provide *distribution transparency*.
- However, trade-offs must be made to have transparency, leading to various techniques to make middleware adaptive.

Distributed Architectural Styles

- **#1 Layered Architectures**

- Basic idea is simple: components are organized in a layered fashion
- Component at layer N is allowed to call components at underlying layer $N - 1$ (but not vice versa)
- This is shown in the diagram
- This model has been widely adopted by the networking community
- A key observation is that control generally flows from layer to layer
- E.g. requests go down the hierarchy whereas the results flow upward.




Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

17

Distributed Architectural Styles (/2)

- **#2 Object-Based Architectures**

- A far looser organization is followed in object-based architectures,
- Each object corresponds to what we have defined as a component,
- These components are connected through a *(remote) procedure call* mechanism.
- This software architecture matches the client-server system architecture (described below).
- Layered & object-based architectures still form the most important styles for large s/w systems




Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

18

Distributed Architectural Styles (/3)

• #3 Event-Based Architectures

- Here, processes communicate thro event propagation, optionally also carrying data
- For distributed systems, event propagation has generally been associated with what are known as *publish/subscribe*
- Idea: processes publish events & m/w ensures that only processes subscribed to the events receive them.
- The main advantage of such systems is that processes are *loosely coupled*.
- Needn't refer to each other explicitly.
- This is also referred to as being *decoupled in space*, or *referentially decoupled*.




Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

19

Distributed Architectural Styles (/4)

• #4 Shared Data-Space Architectures

- Event-based architectures can be combined w data-centered architectures
- Gives what is also known as *shared data spaces*.
- *Essence*: processes now also *decoupled in time*
- Thus need not both be active when communication takes place.
- Also, many shared data spaces use a SQL-like interface to shared repository.
- Means data can be accessed using a description rather than an explicit ref, as per files.




Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

20

SECTION 3.3: SYSTEM ARCHITECTURES: CENTRALIZED & DECENTRALIZED ARCHITECTURES

System Architectures: Centralized Architectures

- *Basic Client–Server Model Characteristics*
 - There are processes offering services (*servers*)
 - There are processes that use services (*clients*)
 - Clients and servers can be on different machines
 - Sometimes Clients can be servers & vice versa
 - Clients follow request/reply model with respect to using services
 - Thinking in terms of Clients requesting Services from Servers aids in the understanding of Distributed Systems



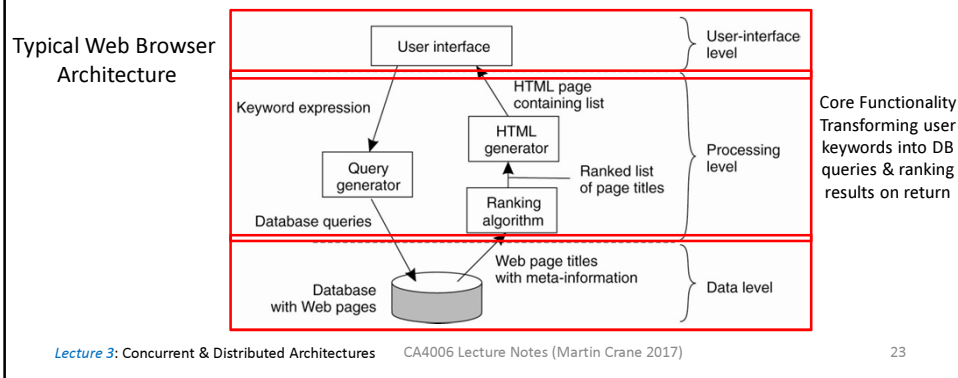
Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

22

System Architectures (/2):

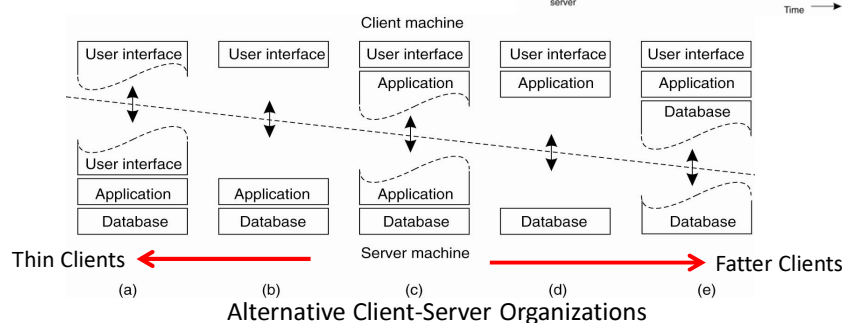
- *Application Layering: Traditional three-layered view*
 1. *User-interface layer* contains units for an application's user interface
 2. *Processing layer* contains the functions of an application, i.e. no specific data
 3. *Data layer* contains data client wants to process thro application components

– Found in many distributed info systems, using traditional DB technology and accompanying applications.



System Architectures: Centralized Architectures (/3)

- *Multi-Tiered Architectures: Variations on traditional 3-layered view*
 1. *Single-tiered*: dumb terminal/mainframe configuration
 2. *Two-tiered*: client/single server configuration
 3. *Three-tiered*: each layer on separate machine (server may act as client)



System Architectures: Decentralized Architectures

- In multi-tiered architectures, the different tiers correspond directly to logical organization of applications – called *Vertical distribution*
- In *horizontal distribution* Client or Server may be split into logically equivalent parts each with own part of data set
- In the last couple of years there has been a tremendous growth in such *peer-to-peer (P2P)* systems:
 - *Structured P2P*: nodes are organized following a specific distributed data structure (usually a Distributed Hash Table)
 - *Unstructured P2P*: nodes have randomly selected neighbours. Each node has a list of neighbours which is constructed in a random way.
 - *Hybrid P2P*: some nodes are appointed special functions in a well-organized fashion

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

25

Decentralized Architectures (/2): Structured P2P Systems

- In virtually all cases, have *overlay networks*
 - This is n/w where nodes are processes & links are communication channels
 - Data is routed over connections setup between nodes.
- As processes can't communicate directly with others, available communication channel must be used (a.k.a. *Application-level Multicasting*)
 - ALM is offered by middleware (in contrast to low-level TCP/IP Multicasting)
 - Basic idea is to organize nodes in a structured overlay n/w such as a logical ring.
 - Specific nodes are made responsible for services based only on their ID.
 - Random key is assigned to a data item from a large (eg 128 bit) identifier space
 - The system provides an operation *LOOKUP(key)* that will efficiently route the lookup request to the associated node.
 - When the key is returned, the network address of node responsible for the data (known as the *successor*) item stored is returned.


Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

26

Decentralized Architectures (/3): Structured P2P Systems: Chord Case Study

• Details of Chord Algorithm

1. Assign random key (*m-bit identifier*) to data item & random number (*m-bit identifier*) to node in system,
2. Implement an efficient & deterministic system to map a data item to a node based on some distance metric,
3. This means that data item should physically be as close to node as possible
4. $LOOKUP(key) \equiv$ returning network address of node responsible for that data item,
5. Do this by routing a request for the data item to responsible node (*successor*).
6. Node with key k falls under the jurisdiction of node with smallest $id \geq k$
7. This process of looking up node's name (& any info stored there) called *name resolution*



Decentralized Architectures (/4): Structured P2P Systems: Chord Case Study


• Building Finger Tables in Chord

Some calculations for Finger tables in the diagram:

$$\begin{aligned} FT_1[1] &= succ(1 + 2^0) = succ(2) = 4 \\ FT_1[2] &= succ(1 + 2^1) = succ(3) = 4 \\ FT_1[3] &= succ(1 + 2^2) = succ(5) = 7 \\ FT_1[4] &= succ(1 + 2^3) = succ(9) = 12 \end{aligned}$$

$$\begin{aligned} FT_4[1] &= succ(4 + 2^0) = succ(5) = 7 \\ FT_4[2] &= succ(4 + 2^1) = succ(6) = 7 \\ FT_4[3] &= succ(4 + 2^2) = succ(8) = 12 \\ FT_4[4] &= succ(4 + 2^3) = succ(12) = 12 \end{aligned}$$

$$\begin{aligned} FT_{15}[1] &= succ(15 + 2^0) = succ(16) = succ(0) = 1 \\ FT_{15}[2] &= succ(15 + 2^1) = succ(17) = succ(1) = 1 \\ FT_{15}[3] &= succ(15 + 2^2) = succ(19) = succ(3) = 4 \\ FT_{15}[4] &= succ(15 + 2^3) = succ(23) = succ(7) = 7 \end{aligned}$$




Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2015)

29

Decentralized Architectures (/5): Structured P2P Systems: Chord Case Study

• Principle of Joining a System in Chord

- Node wanting to join system starts by generating random identifier $id = 26$.
 - Then node simply contacts an arbitrary node & does a lookup on id ,
 - Returns address of $succ(id) = 32$, node responsible for looking after id
 - Next, node simply contacts $succ(id)$ & its predecessor & inserts self in ring
 - This consists of updating the finger tables.
 - Insertion also yields that each data item whose key is now associated with node id , is transferred from $succ(id)$.
- Chord scheme requires that each node also stores info on its predecessor.



Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

30

Decentralized Architectures (/6): Structured P2P Systems: Chord Case Study

- *Problems in Chord*
- Logical organization of overlay nodes may lead to erratic msg transfers in underlying Internet: node k , node $\text{succ}(k)$ may be far apart.
 - *Topology-aware node assignment:*
 - When assigning an ID to a node, make sure that nodes close in the ID space are also close in the network.
 - Can be very difficult.
 - *Proximity routing:*
 - Maintain more than one possible successor, and forward to the closest.
 - Example: in Chord $FT_p[i]$ points to first node in $[p + 2^{i-1}, p + 2^i - 1]$.
 - Node p can also store pointers to other nodes in the interval.
 - *Proximity neighbour selection:*
 - When there is a choice of selecting who your neighbour will be (not in Chord), pick the closest one.

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

31

Decentralized Architectures (/7): Structured P2P Systems: Pastry Case Study

- *Properties of Pastry:*
- **PASTRY** is an implementation of a Distributed Hash Table (DHT) algorithm for P2P routing overlay
- Salient features:
 - Fully decentralized
 - Scalable
 - High fault tolerance
- Each node is identified by a unique 128 bit node id (*NodeId*) generated randomly so each has same probability of being chosen
- Node with similar *NodeId* may be geographically far apart
- Given a *key*, PASTRY can deliver a message to node with closest *NodeId* to *key* within $\log_{2^b} N$ steps,
 - where b is a configuration parameter (usually $b = 4$)
 - and N is the number of nodes

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

32

Decentralized Architectures (/8): Structured P2P Systems: Pastry Case Study (/2)

- *Pastry Routing Algorithm:*
- Given want to find PASTRY n/w node with *NodeId* closest to given *key*
 - Note that *NodeId* & *key* are both 128 bit sequences
 - Both *NodeId* & *key* can be thought as sequence of digits with base 2^b
- *Routing idea:*
 1. Each routing step, node normally forwards message to a node whose *NodeId* shares with *key* a prefix min. 1 digit longer than *key* shares with present node.
 2. If such a node unknown, message is forwarded to a node that shares same prefix of actual node but its *NodeId* is numerically closer to *key*

Decentralized Architectures (/9): Structured P2P Systems: Pastry Case Study (/3)

- *State of a Node in Pastry:*
- Each PASTRY node has a *state* consisting of:
 - A *routing table* *R*
 - used in the first phase of the routing (*long distances*)
 - A *neighbourhood set* *M*
 - contains *NodeId* & IP addresses of the $|M|$ nodes which are *closest* (according to a *metric*, e.g. geog. or ping distance) to considered node
 - A *leaf set* *L*
 - contains *NodeId* & IP addresses of $|L|/2$ nodes with *NodeId* *numerically closest on smaller* side of present *NodeId*,
 - and $|L|/2$ nodes with *NodeId* *numerically closest on the larger* side of present *NodeId*.
 - *L* usually taken to be 16

Decentralized Architectures (/10): Structured P2P Systems: Pastry Case Study (/4)

- *Routing table in Pastry:*
- This is a $\lceil \log_{2^b} N \rceil$ rows $\times (2^b - 1)$ columns table
 - where $\lceil \log_{2^b} N \rceil$ is the max number of hops between any pair of nodes
 - b is the configuration parameter (usually 4) and
 - N is the number of PASTRY nodes in the network
- The $2^b - 1$ entries at row n each refer to a node whose *NodeID* shares the present node *NodeID* in the first n digits
- However the $(n + 1)$ th digit has one of the $2^b - 1$ possible values other than $(n + 1)$ th digit digit in the present node *id*.
- The choice of b is a choice between the size of the populated part of the Routing table ($\lceil \log_{2^b} N \rceil \times (2^b - 1)$ entries) & max number of hops.
 - e.g. a value of $b = 4$ and $N = 10^6$ nodes gives ~ 75 entries and ~ 5 hops
 - while $b = 4$ and $N = 10^9$ Nodes gives ~ 105 entries and ~ 7 hops

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

35

Decentralized Architectures (/11): Structured P2P Systems: Pastry Case Study (/5)

- *Example Routing Table R in Pastry:*
- $N = 1024$ Nodes, $b = 2$ so $\lceil \log_{2^b} N \rceil = 5$ rows, $2^b - 1 = 3$ columns
 - Row i : Holds ids of Nodes whose IDs share an i digit prefix with Node
 - Column j : digit $(i + 1) = j$
 - Contains topologically closest node that meets these criteria

Shared prefix length with NodeID

Digit at position $i+1$

i	j	0	1	2	3
0		01230	13320	22222	
1		30331	31230	-	33123
2		32012	-	32212	32301
3		-	32110	32121	32131
4		32100	-	32102	32103

These entries match node 32101's ID

Topologically closest with prefix length i & digit $(i+1)=j$

Possible node 33xyz
33123 is topologically closest node

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

36

Decentralized Architectures (/12): Structured P2P Systems: Pastry Case Study (/6)

- *Example Routing in Pastry: $N = 1024$ Nodes, $b = 2, L = 8$*
- Leaf Table L for NodeID 32101
 - $L/2$ smaller, $L/2$ larger
 - Fixed maximum size
 - Similar to Chord's finger table
 - Used for routing and recovery from departures of nodes
- Neighbour Set M
 - Contains nearby nodes (based on some scalar proximity metric e.g. geography, latency, IP hops etc)
 - Fixed maximum size
 - Irrelevant for routing

Smaller NodeID's		Larger NodeID's	
32100	32023	32110	32121
32012	32022	32123	32120

Lecture 3: Concurrent & Distributed Architectures

CA4006 Lecture Notes (Martin Crane 2017)

37

Decentralized Architectures (/13): Structured P2P Systems: Pastry Case Study (/7)

- *Routing Algorithm of Packet with NodeID A , key D (both 128 bit):*
 - (1) if $(L_{-|L|/2} \leq D \leq L_{|L|/2})$ then
 - (2) *// D is in the Leaf Node Set*
 - (3) forward to L_i , such that $|D - L_i|$ is minimal, i.e. closest NodeID in L
 - (4) else
 - (5) *// search for a node with longer shared prefix in the routing table*
 - (6) Let $l = shl(D, A)$
 - (7) if $(R_l^{D_l} \neq null)$ then
 - (8) forward to $R_l^{D_l}$ *// entry in routing table row l , column D_l*
 - (9) *D_l is the value of the l 's digit in the key D*
 - (10) else
 - (11) *// rarely*
 - (12) forward to $T \in L \cup R \cup M$ such that
 - (13) $shl(T, D) \geq l, |T - D| < |A - D|$
 - (14) search for node T with longest prefix out of merged set

Lecture 3: Concurrent & Distributed Architectures

CA4006 Lecture Notes (Martin Crane 2017)

38

Decentralized Architectures (/14): Unstructured P2P Systems

- Many unstructured P2P systems try to maintain a random graph
- Basic principle is for each node is required to contact a randomly selected other node:
 - Let each peer maintain a partial view of the network, consisting of c other nodes
 - Each node P periodically selects a node Q from its partial view
 - P and Q exchange information and exchange members from their respective partial views
- It turns out that, depending on the exchange, randomness, but also robustness of the network can be maintained.


Lecture 3: Concurrent & Distributed Architectures

CA4006 Lecture Notes (Martin Crane 2017)

39

Decentralized Architectures (/15): Unstructured P2P Systems (/2)

- *Topology Management of Overlay Networks*
- Basic idea is to distinguish two layers:
 1. maintain random partial views in lowest layer;
 2. be selective on who you keep in higher-layer partial view.



- Lower layer feeds upper layer with random nodes; upper layer is selective when keeping references (e.g. based on distance).

Lecture 3: Concurrent & Distributed Architectures

CA4006 Lecture Notes (Martin Crane 2017)

40


Decentralized Architectures (/16): Unstructured P2P Systems (/3)

- *Topology Management of Overlay Networks (cont'd)*
- To construct a torus, Consider a $N \times N$ grid.

Keep only refs to nearest neighbours:

$$\|(a_1, a_2) - (b_1, b_2)\| = d_1 + d_2$$

$$d_i = \min\{N - |a_i - b_i|, |a_i - b_i|\}$$



Time →

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

41


Decentralized Architectures (/17): Unstructured P2P Systems (/4)

- *Topology Management of Overlay Networks (cont'd)*
- To construct a torus, Consider a $N \times N$ grid.

Keep only refs to nearest neighbours:

$$\|(a_1, a_2) - (b_1, b_2)\| = d_1 + d_2$$

$$d_i = \min\{N - |a_i - b_i|, |a_i - b_i|\}$$



Here, there are two points:

$(a_1, a_2) = (1,3)$ and $(b_1, b_2) = (3,0)$

hence

$d_1 = \min\{4 - 2, 2\} = 2$ (both paths same length)

and


$d_2 = \min\{4 - 3, 3\} = 1$ (green path is shorter)

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

42

Decentralized Architectures (/18): Unstructured P2P Systems (/5)

- *Topology Management of Overlay Networks (cont'd)*
- Explanation
 - With minimum distance condition, a toroidal shape emerges.




Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

43

Decentralized Architectures (/19): Hybrid Architectures: C-S combined with P2P

- Example: *Edge-server* architectures, which are often used for *Content Delivery Networks*



Viewing the Internet as consisting of a collection of edge servers.

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

44

Decentralized Architectures (/20): Hybrid Architectures: C-S with P2P (/2)

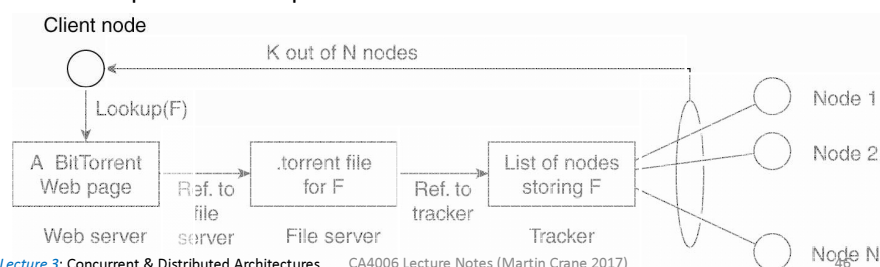
- *Internet as consisting of a collection of edge servers*
- An important class of distributed systems that is organized according to a hybrid architecture is formed by *edge-server systems*.
- Such systems are deployed on the Internet where servers are placed "at the edge" of the network.
 - Edge is formed by boundary between enterprise n/ws and actual Internet, (for example, as provided by an ISP).
 - Likewise, where end users at home connect to the Internet through their ISP, the ISP can be considered as residing at edge of Internet.
 - Edge-Server thus serves content and optimises delivery
- *Content Delivery Networks* offers storage of copies of webpages for rapid reaccessing.

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

45

Decentralized Architectures (/21): Hybrid Architectures: C-S combined with P2P

- *Example: Hybrid Architectures: C/S with P2P – BitTorrent*
- *Basic Idea:* **Tracker** (server with list of active nodes to download chunks of file) gives single copy (*seed*) of file (F), *swarm* is all nodes with some/all of F
- *Steps:*
 1. Client Node does a Lookup on F,
 2. BT webpage gives ref to file server with **.torrent** file for F (with Tracker).
 3. BT Client s/w talks to tracker to find other BT Nodes with whole/part of F.
 4. Tracker identifies swarm (i.e. connected peers sending/receiving) F.
 5. Tracker helps client trade pieces of F needed with others in swarm.



Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

Architecture V Middleware

• *Architecture and Middleware*

- Considering the architectural issues above, a question that comes to mind is where middleware fits in.
- Important aim is to give a degree of distribution transparency, i.e. try to hide data distribution, processing, and control from applications.
- What is commonly seen in practice is that middleware systems actually follow a specific architectural style.
- The chosen style may not be optimal in all cases.
- So may need to (dynamically) adapt behaviour of the middleware.

Interceptors

- These intercept usual flow of control when invoking a remote object.
- Thus they allow other (application specific) code to be executed.
- This is demonstrated in the diagram (over)

Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

47

Architecture V Middleware : Interceptors (2)

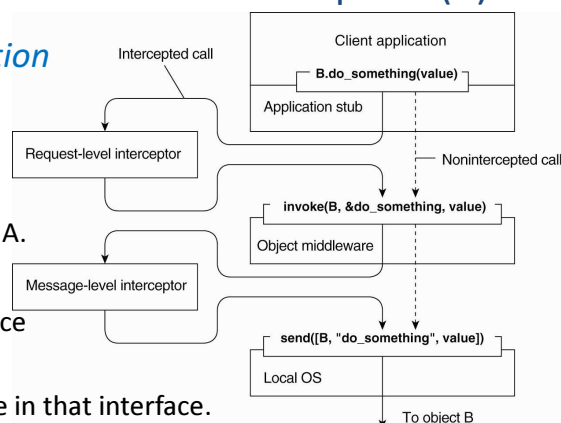
• *Remote Object Invocation*

- Basic idea:

Object A can call a method belonging to object B, living on a different machine to A.

- Steps:

1. A offered a local interface (same as B's).
2. A calls method available in that interface.
3. A's call transformed into a generic object invocation, enabled thro a general object-invocation interface offered by m/w at A's machine.
4. Finally, GOI is transformed into a message sent thro the transport-level network interface offered by A's local operating system.



Lecture 3: Concurrent & Distributed Architectures CA4006 Lecture Notes (Martin Crane 2017)

48

Summary

- Flynn's Taxonomy is a classic but still useful way to classify architectures:
 - SISD, SIMD, MIMD can still be identified in supercomputers today
 - MIMD can be split into Tight & Loose Coupling
- Software Architectures for Distributed Systems divide into:
 - Layered, Object-/Event-based, Shared Dataspace Architectures
- System Architectures
 - Centralized Architectures:
 - 2 & multi-tiered architectures
 - Fat & Thin Clients
 - Decentralized Architectures can be divided into
 - Structured P2P Systems (e.g. Chord & Pastry Routing algorithms)
 - Unstructured P2P Systems
 - Hybrid Systems (e.g. BitTorrent)
- Middleware can sometimes be used to fill in for architecture